

FAQ & Additional Tips: Taiga Project

Table of Contents

- [I. Platform Migration](#)
- [II. General Questions](#)
- [III. Token Counting, Constraints, & the "Turn-to-Turn" Architecture](#)
- [IV. Workspace, Codebases, & Environment Configuration](#)
- [V. Platform Errors, Safety Blocks, & UI Troubleshooting](#)
- [VI. Task Evolution & Evaluation Criteria](#)
- [VII. Administrative, Hubstaff, & Support Protocols](#)

I. Platform Migration

Task Creation & Visibility in Labelbox

Q: How do I create a task?

Create new tasks in the Taiga platform, not in Labelbox. Watch the instructional video for step-by-step guidance: [Submitting a task in Labelbox](#). **DO NOT** click "New" in Labelbox.

Q: When will a task I just created appear in Labelbox?

A task appears in Labelbox only after you complete the first turn (writing rationales) and clicking on "Submit Preference".

Q: How do I view the tasks I've created in Labelbox?

Open the project dashboard in Labelbox and click the "Tasks" tab. All your assigned tasks are listed there.

Task Submission

Q: When my submission is finished do I need to submit in Taiga?

No, when a submission is finished, fill out the submission form in Labelbox and start a new live compare task. **DO NOT** "submit to customer" within Taiga.

Q: Should I submit my task to the client?

No. Never submit the task to the client.

Q: Can I click ?

If you click "Finish & Start New Conversation" you will not be able to continue adding turns. You can click this once the submission has reached 1M tokens indicating completion and the task has been submitted in Labelbox, but this is not required. As long as the preferences are submitted and you have completed the submission form in Labelbox this step is not required.

Q: What do I do if a task isn't finished by the end of Sunday?

Never submit a partial task; submit only once it is complete. If you can't finish by Sunday, submit it during the next pay period. It will be counted toward, and paid in full during that next pay period.

Q: When should I submit my task?

There are only three situations in which you should submit a task. Go to the [Submitting Your Task](#) tab to learn about these situations.

Time Tracking


Q: My time tracked isn't reflected in Labelbox. What should I do?

Don't worry about this; it doesn't affect payment. Payment is based on the time reported when you complete the task. Use the time in Hubstaff as your reference, not the time shown in Labelbox.

Known Platform Issues

Q: What does the feedback on my tasks mean, and why do I have a score of 0.0?

Ignore this. This feature isn't working properly.



II. General Questions

Q: Can I use Grammarly or other AI-assistive tools to improve my writing?

No.

Q: When I finish a turn, what do I do if I haven't hit the low-context-error yet?

You need to continue the conversation by clicking "Submit preference". Do NOT press "Finish & Start New Conversation".

If you selected a preference other than the smallest preference towards a response, which individual axes held the most weight in your overall preference selection? Please list up to 2. (optional)

Select axes (0) ▾

If you selected the smallest preference, what was the quality of both models (optional)

Select quality (0) ▾

Select at least 1

Choose the response that is better overall.* ⓪


A	A	A	A	B	B	B	B
---	---	---	---	---	---	---	---

A is better B is better

Submit & start new Submit preference

Q: I have closed the Taiga window in my browser, how do I get back into the submission I was working on?

From the main Taiga screen click on "Live Compare". Instead of entering a new prompt, hover over the left side of the screen and click on the submission you were working on.



✓ **Q: I got an error "Parent run has no container checkpoint". Can I proceed with my task?**

No. The checkpoint was not created. You can diagnose this early. After each turn, you should see a button "Save Final State" after the transcript. If you don't, something likely went wrong with the task, and you should submit the task through your Labelbox dashboard, following the instructions for [Handling Platform Errors](#).

Q: Does the model have internet access?

No, but the models do have access to client maintained package managers. These currently include pip, npm, and uv.

Q: How do I get the LiveCompare-ID for submission?

Navigate to Environment > Jobs > "Show Only My Jobs"



The format should be live-compare-XXXXXX. In the above example, the live-compare-ID would be "live-compare-bb677646". Don't include the repeating "live-compare" or the "__cont..".

Q: Is there a time limit to starting and finishing a task?

Yes, a task must be started and finished within 7 days, otherwise the checkpoint will be invalidated, and you can't continue working on that task. This is a limit set by the client; we do not have any control over this.

III. Token Counting, Constraints, & the "Turn-to-Turn" Architecture

Q: How is the "1 million token" target tracked?

The client that sends requests to the server is the context for each task. It includes the previous 4 million

The signal that you've reached the target is the context-low error itself, which fires near 1 million tokens and reports the exact final token count you submit. Don't rely on the live "tokens remaining" counter as a precise gauge.

Q: Why does the "tokens remaining" counter sometimes change unexpectedly between turns?

In rare cases the counter can appear to jump up or down between turns, due to nuances in the platform's orchestration layer (such as sub-agent context drops or variations in prompt evaluation). Don't read too much into any single turn's number or watch it too closely; keep pushing the conversation forward and let the context-low error tell you when you've hit the limit.

Q: How many turns should a task take? What if I hit the 30-turn maximum before 1 million tokens?

Most reach the 1 million token target in 10 to 20. If you hit 30 turns, stop and submit the task immediately. If you notice your conversation isn't consuming significant tokens, adjust your strategy to generate heavier text and code payloads naturally.

IV. Workspace, Codebases, & Environment Configuration

Q: Can I upload additional source files or patches mid-conversation?

Yes.

Q: Can I push the task codebase to my own GitHub repo for reference?

No. Pushing the codebase elsewhere violates the client's security policy. Everything you create within the platform must stay there.

✓ **Q: What are the strict specifications for uploading a codebase on Turn 1?**

- Aim to package everything the model needs into the initial Turn 1 upload. Front-loading keeps the setup clean.
- The size of each uploaded file must strictly remain under a 50MB cap.
- You must compress your directory as a .zip file. Tarballs (.tar, .tar.gz) are poorly supported and frequently fail to unpack correctly within the harness.
- Avoid simple "toy" repos, automated boilerplates, standard code templates, or basic classroom/bootcamp demo projects. The platform requires codebases that reflect real-world use.

Q: The repository I want to use is larger than the 50MB platform cap. How can I shrink it?

Large repositories are usually weighed down by binary objects or local modules. You can safely delete node_modules folders, heavy build assets, compiled binaries, local images, or raw videos.

Q: Can the model install system-wide dependencies?

The environment is an unprivileged sandbox container. The model does not have global sudo permissions; commands like sudo apt-get install will fail out of the box. However, highly capable models may attempt to download Debian packages, extract them locally, and map paths into user space to circumvent directory permission barriers. This behavior is technically a clever routing mechanism and is acceptable to document.

Q: Which programming languages are supported?

- Python — the most supported language in the sandbox; native and fully functional.
- JavaScript / TypeScript / ReactJS — fully supported and recommended.
- C / C++ — supported and viable.
- Java — not recommended. Support is inconsistent due to JDK, Maven, and Gradle compatibility problems, and it can cause system instability or potential data loss. Avoid building DevOps/CI-CD tasks around Java for this reason.
- PHP — does not function in the environment, specifically when used with a Laravel back end.
- MATLAB — no specific prohibition, but unverified by the team. You may attempt it to determine viability at your own discretion.
- Ruby — not supported.
- Rust — partially available via Cargo, but lacks mature, localized compilers in the sandbox.

Always task using a language you are highly proficient in so you can accurately audit behavioral flaws.

Q: Can I use public datasets such as those from Kaggle?

Yes, as long as the data is genuinely public and not private or sensitive. Confirm a dataset's public status before using it, and stay within the 50MB upload cap using ZIP compression.

Q: Can I reuse a codebase across tasks, and how should the model analyze it?

Reusing the same codebase is allowed up to 5 times, after which you should rotate to a fresh source. When analyzing, prefer partial batches over feeding the whole codebase at once; smaller batches often produce better model performance and help avoid the token-budget cut-offs where models truncate large reviews.

Q: If I use a personal repository to test models on the platform, can I still publish my app later? Does it violate the NDA?

Yes, you can still publish your original app. While any downstream modifications, alternative variants, or direct code branches generated by the platform belong exclusively to the client under the NDA, your original, pre-existing source code remains entirely your intellectual property.

V. Platform Errors, Safety Blocks, & UI Troubleshooting

Action Step: If you were booted due to a background VPN issue, reach out directly to Marcos on Discord. For general access blocking, fill out the official Access Issues Form and send a DM to the project leads (Wesley or Amicable-Lead).

Note: Not every removal is a penalty. Some recent cases were reassignments to a different project under an unusually tight client timeline, not disciplinary action.

Q: What does an "overload" or "timeout" error mean, and how do I fix it?

These are standard environment stability issues. Do not panic or repeatedly mash buttons. Simply stop interacting with the interface and wait for the models to regenerate their response.

Q: The platform went down while I was in the middle of a task. Can I still write and save my rationales?

No. If an outage occurs, freeze all interaction with the platform. Writing your rationales during an outage introduces an unnecessary risk of data loss. You may draft your notes locally in an external text file if you wish, but do not submit anything until a team lead explicitly announces on Discord that the platform is stable.

Q: What is the diagnostic checklist if a task freezes or throws a checkpoint error?

Look at the bottom of the interface for a "Save Final State" button. If that button is missing and your execution logs show a failed checkpoint creation or a failed checkpoint stack trace, the task state is fundamentally broken. If you are in a later turn (e.g., turn 2, 3, etc.), submit the task in labelbox, mark "no" for completion due to error and explain the error, then start a new submission in Taiga. You are fully permitted to spin up a brand new task using the exact same Turn 1 generator prompt to rerun the workflow if the error is checkpoint related and in turn 1.

▼ **Q: How do I handle an explicit "Effort Configuration Blocks" error?**

This error occurs during turn submission and preference syncing. If a hard refresh or clearing your browser cache/cookies fails to resolve it, do not keep spamming the button. Stop, copy your Live Compare ID and environment URL, and report it immediately via the Incident Tracker. If you are in a later turn (e.g., turn 2, 3, etc.) and unable to continue the submission due to the error, submit the task in labelbox, mark "no" for completion due to error and explain the error, then start a new submission in Taiga.

Q: My task was aborted by an "Anthropic Trust & Safety Safeguard Block". What should I do?

This block triggers an empty response when a prompt inadvertently fires the client's automated policy classifier. This is common when testing security configurations, penetration testing scripts, or data parsers. If it happens on Turn 1, it is an immovable false positive; abandon the task. If you are in a later turn (e.g., turn 2, 3, etc.), submit the task in labelbox, mark "no" for completion due to error and explain the error, then start a new submission in Taiga. Please do not attempt cybersecurity or biology prompts.

Q: Why does the system crash when I paste a link or include an image attachment?

The sandbox completely lacks an image-processing layer. Embedding images, screenshot URLs, or SVG data URLs will instantly stall the sampling subprocess. Do not use visual attachments under any circumstances. Translate any visual context into raw descriptive text directly inside your prompt.

▼ **Q: I am experiencing an annoying double-paste glitch inside the rationale fields. Is there a fix?**

This is a known UI input bug affecting the reasoning text blocks. When copying and pasting text or moving a highlighted citation, it may duplicate. Review your text boxes carefully before committing to clean up duplicate lines.

Q: The Code Diff feature disappeared or is showing bugs, and how do I download model outputs?

The built-in code diff engine can be highly unreliable, especially on long-running tasks. Do not rely on it; if it fails, manually download the relevant files/artifacts from the model to inspect changes. There is no single button to download all outputs at once, but if you set your output path to /temp/outputs in settings you can download at the end of each turn — and transcripts remain downloadable even if you didn't configure this beforehand.

Q: Is it possible to unselect a rating axis component once I have clicked it?

Currently, the platform interface does not allow you to unselect a behavioral axis rating once an option has been clicked. If you click one accidentally, you must complete it and defend it.

Q: Can I use multiple monitors while tasking?

Yes. Using multiple monitors will not conflict with the platform or your environment tracking. Hubstaff safely records activity across multiple displays.

VI. Task Evolution & Evaluation Criteria

Q: How do I approach Deployment, Maintenance, and CI/CD categories given platform limits?

Because the sandbox lacks container access and can't actually run real deployments, you act as the manual operator. Frame the conversation so the model generates configurations (e.g., Terraform, Kubernetes manifests), setup scripts, or CI/CD pipelines and metrics, while you provide output results or mock terminal errors back to the model. It is legitimate to have the model create and instantiate these components even though the sandbox can't actually execute them. Avoid asking it to run deployment commands agentially.

VII. Administrative, Hubstaff, & Support Protocols

Q: When does the contractual work week reset?

The tracking week resets at 00:00 UTC (the Sunday-to-Monday boundary). No formal time-off requests are required if you will miss hours, as it is a flexible rolling basis.

Q: Do I keep Hubstaff running while waiting for models to finish "Sampling"?

Track only active working time. While you're actively auditing the model's output as it streams, keep Hubstaff running. But pause Hubstaff during long-running sampling sessions where you're just waiting and not working, and resume as soon as you're back to active work. Also pause your timer (and report it) if the platform crashes, logs out, or enters an infinite server-side loading hang.

Q: Labelbox states a "14-hour max payout limit" per task. Will I get paid if a task takes longer?

Yes. The 14-hour language on Labelbox is a default setting and does not reflect actual backend payouts. Because these are deep, exhaustive tasks, going over 14 hours is understandable in many cases. Payouts are directly calculated based on your accumulated Hubstaff hours. Ensure you are strictly following the expected Average Handle Time (AHT) of 60 to 100 minutes per individual turn.

Q: How do I report my hours if Hubstaff resets at midnight?

Go into [Hubstaff](#) to see hours.

Q: Do office hours count as project time on Hubstaff?

No. Office hours are not considered part of project tasks and should not be tracked on Hubstaff — only your actual tasking time counts.

Q: How long will the project run?

There's no fixed end date; Longevity depends on the quality of data contributions. If you're qualified, you're free to move between projects on the platform.